

**TERM: L3-AII & L3-ELNI**

**SEMESTER: 5**

**AY: 2023-2024**

# Abdelbacet Mhamdi

Dr.-Ing. in Electrical Engineering

Senior Lecturer at ISET Bizerte

abdelbacet.mhamdi@bizerte.r-iset.tn

## MACHINE LEARNING

LAB MANUAL



**Institut Supérieur des Études Technologiques de Bizerte**

---

Available @ <https://github.com/a-mhamdi/mlpy/>



# --- HONOR CODE ---

THE UNIVERSITY OF NORTH CAROLINA AT CHAPEL HILL

Department of Physics and Astronomy

<http://physics.unc.edu/undergraduate-program/labs/general-info/>

“During this course, you will be working with one or more partners with whom you may discuss any points concerning laboratory work. However, you must write your lab report, in your own words.

Lab reports that contain identical language are not acceptable, so do not copy your lab partner’s writing.

If there is a problem with your data, include an explanation in your report. Recognition of a mistake and a well-reasoned explanation is more important than having high-quality data, and will be rewarded accordingly by your instructor. A lab report containing data that is inconsistent with the original data sheet will be considered a violation of the Honor Code.

Falsification of data or plagiarism of a report will result in prosecution of the offender(s) under the University Honor Code.

On your first lab report you must write out the entire honor pledge:

---

**The work presented in this report is my own, and the data was obtained by my lab partner and me during the lab period.**

---




On future reports, you may simply write “Laboratory Honor Pledge” and sign your name.”

# Contents

<b>1</b>	<b><i>Python</i> Onramp</b>	<b>1</b>
<b>2</b>	<b>Linear Regression</b>	<b>7</b>
<b>3</b>	<b><math>k</math>-NN for Classification</b>	<b>14</b>
<b>4</b>	<b>K-Means for Clustering</b>	<b>19</b>
<b>5</b>	<b>Binary Classifier using ANN</b>	<b>25</b>


---

In order to activate the virtual environment and launch **Jupyter Notebook**, we recommend you to proceed as follow

- ① Press simultaneously the keys  &  on the keyboard. This will open the dialog box **Run**;
- ② Then enter `cmd` in the command line and confirm with  key on the keyboard;
- ③ Type the instruction `mlpy.bat` in the console prompt line;



```
Command Prompt
C:\Users\admin> mlpy.bat
```

- ④ Finally press the  key.

---

**LEAVE THE SYSTEM CONSOLE ACTIVE.**

# 1 | Python Onramp

Student's name	.....	.....	.....
	.....	.....	.....
	.....	.....	.....
Score /20	.....	.....	.....

## Detailed Credits

Anticipation (4 points)	.....	.....	.....
Management (2 points)	.....	.....	.....
Testing (7 points)	.....	.....	.....
Data Logging (3 points)	.....	.....	.....
Interpretation (4 points)	.....	.....	.....

### Motivations

- ★ *Python* is a popular programming language in the field of machine learning because it is relatively easy to learn and has a wide range of libraries and frameworks that support machine learning tasks.
- ★ *Python* has a large and active community of developers, which means that there are many resources available online, such as tutorials, documentation, and online forums, to help us learn and troubleshoot our code.
- ★ Many machine learning tools and frameworks, such as *TensorFlow* and *scikit-learn*, are written in *Python*, which makes it easy to integrate these tools into *Python* programs.
- ★ *Python* is a versatile language that can be used for a wide range of applications beyond machine learning, including web development, data analysis, and scientific computing. Learning *Python* can therefore open up many career opportunities for us.



The notebook is available at <https://github.com/a-mhamdi/mlpy/> → Codes → Python → py-onramp.ipynb

# k-means-clustering

November 30, 2023

## 1 Machine Learning

Textbook is available @ <https://www.github.com/a-mhamdi/mlpy>

---

In unsupervised learning, the algorithm is given a dataset and is asked to learn the underlying structure of the data. The goal is to find patterns or relationships in the data that can be used to group the data points into clusters or to reduce the dimensionality of the data.

Some examples of unsupervised learning algorithms include: 1.  $K$ -means clustering; 1. Principal Component Analysis (PCA); and 1. Autoencoders.

These algorithms can be used for tasks such as image compression, anomaly detection, and customer segmentation.

### 1.1 K-Means Clustering

$K$ -means clustering is a method of unsupervised machine learning used to partition a dataset into  $k$  clusters, where  $k$  is a user-specified number. The goal of  $K$ -means clustering is to minimize the sum of squared distances between the points in each cluster and its centroid.

#### 1.1.1 Importing the libraries

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
```

```
[2]: # Show plots in an interactive format, e.g., zooming, saving, etc
%matplotlib inline
```

```
[3]: plt.style.use('ggplot')
```

#### 1.1.2 Importing the dataset

```
[4]: df = pd.read_csv('./datasets/Mall_Customers.csv')
```

```
[5]: df.head()
```

```
[5]:
```

	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	Male	19	15	39
1	2	Male	21	15	81
2	3	Female	20	16	6
3	4	Female	23	16	77
4	5	Female	31	17	40

```
[6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   CustomerID            200 non-null   int64
1   Gender                200 non-null   object
2   Age                   200 non-null   int64
3   Annual Income (k$)    200 non-null   int64
4   Spending Score (1-100) 200 non-null   int64
dtypes: int64(4), object(1)
memory usage: 7.9+ KB
```

```
[7]: df.describe()
```

```
[7]:
```

	CustomerID	Age	Annual Income (k\$)	Spending Score (1-100)
count	200.000000	200.000000	200.000000	200.000000
mean	100.500000	38.850000	60.560000	50.200000
std	57.879185	13.969007	26.264721	25.823522
min	1.000000	18.000000	15.000000	1.000000
25%	50.750000	28.750000	41.500000	34.750000
50%	100.500000	36.000000	61.500000	50.000000
75%	150.250000	49.000000	78.000000	73.000000
max	200.000000	70.000000	137.000000	99.000000

```
[8]: df.rename(columns={'Annual Income (k$)': 'Annual Income', 'Spending Score_
↳ (1-100)': 'Spending Score'}, inplace=True)
```

```
[9]: X = df.drop(columns=['CustomerID', 'Age', 'Gender']).values
X[:10, :]
```

```
[9]: array([[15, 39],
           [15, 81],
           [16,  6],
           [16, 77],
           [17, 40],
           [17, 76],
           [18,  6],
```



```
[18, 94],
[19, 3],
[19, 72]])
```

Import KMeans class

```
[10]: from sklearn.cluster import KMeans
```

**OPTIONAL: IF NOT FAMILIAR WITH KMEANS, FEEL FREE TO SKIP THE FOLLOWING CELL**

---

Using the elbow method to find the optimal number of clusters

```
[11]: """
wcss = []

for i in range(1, 11):
    kmeans = KMeans(n_clusters=i,
                    init='k-means++', # Init method
                    random_state=123) # Random seed for reproducibility
    kmeans.fit(X)
    wcss.append(kmeans.inertia_)

plt.plot(range(1, 11), wcss)
plt.title('The Elbow Method')
plt.xlabel('Number of clusters')
plt.ylabel('WCSS')
plt.grid()
"""
```

```
[11]: "\nwcss = []\n\nfor i in range(1, 11):\n    kmeans = KMeans(n_clusters=i,\n    init='k-means++', # Init method\n                                random_state=123) # Random\n    seed for reproducibility\n    kmeans.fit(X)\n    wcss.append(kmeans.inertia_)\n\nplt.plot(range(1, 11), wcss)\nplt.title('The\nElbow Method')\nplt.xlabel('Number of\nclusters')\nplt.ylabel('WCSS')\nplt.grid()\n"
```

---

### 1.1.3 Training the K-Means model on the dataset

This code will create a *K*-means model with 5 clusters and fit it to the data. It will then make predictions about which cluster each data point belongs to

```
[12]: kmeans = KMeans(n_clusters=5, init='k-means++', random_state=123)
y_pred = kmeans.fit_predict(X)
```

```
/home/mhamdi/MEGA/git-repos/AI-ML-  
DL/mlpy/Codes/Python/pvenv/lib/python3.10/site-  
packages/sklearn/cluster/_kmeans.py:870: FutureWarning: The default value of  
`n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init`  
explicitly to suppress the warning  
warnings.warn(  


```

```
[13]: centers = kmeans.cluster_centers_  
centers
```

```
[13]: array([[55.2962963 , 49.51851852],  
            [86.53846154, 82.12820513],  
            [88.2       , 17.11428571],  
            [26.30434783, 20.91304348],  
            [25.72727273, 79.36363636]])
```

#### 1.1.4 Visualizing the clusters

```
[14]: fig, ax = plt.subplots()  
scatter = ax.scatter(X[:, 0], X[:, 1], c=y_pred, s=100)  
legend = ax.legend(*scatter.legend_elements(), title='Clusters')  
ax.add_artist(legend)  
ax.scatter(centers[:, 0], centers[:, 1], c='red', s=200)  
ax.set_title('Clusters of customers')  
ax.set_xlabel('Annual Income')  
ax.set_ylabel('Spending Score')  
ax.grid()
```



Unsupervised learning can be useful when there is no labeled training data available, or when the goal is to discover patterns or relationships in the data rather than to make predictions. However, it can be more difficult to evaluate the performance of unsupervised learning algorithms, as there is no ground truth to compare the predictions to.

*K*-means clustering is a fast and efficient method for clustering large datasets, and is often used as a baseline method for comparison with other clustering algorithms. However, it can be sensitive to the initial selection of centroids, and may not always find the optimal clusters if the data is not well-separated or has a non-convex shape. It is also limited to spherical clusters and may not work well for clusters with more complex shapes.

## 2 | Linear Regression

<b>Student's name</b>	.....	.....	.....
	.....	.....	.....
	.....	.....	.....
<b>Score</b> <b>/20</b>	.....	.....	.....

### Detailed Credits

<b>Anticipation (4 points)</b>	.....	.....	.....
<b>Management (2 points)</b>	.....	.....	.....
<b>Testing (7 points)</b>	.....	.....	.....
<b>Data Logging (3 points)</b>	.....	.....	.....
<b>Interpretation (4 points)</b>	.....	.....	.....

### Motivations

- ★ Linear regression is a fundamental statistical technique that is widely used in many fields, including economics, finance, biology, and computer science. It is a simple and effective way to model the relationship between a dependent variable and one or more independent variables.
- ★ Linear regression is relatively easy to understand and implement, making it a good starting point for us who are new to statistical modeling. It is also a good foundation for learning more advanced statistical techniques, such as multivariate or logistic regression.
- ★ Linear regression can be an useful tool for making predictions and understanding the underlying trends in data. It can help us to better understand and analyze data, and to make informed decisions based on our findings.



The notebook is available at <https://github.com/a-mhamdi/mlpy/> → Codes → Python → multiple-linear-regression.ipynb

# multiple-linear-regression

November 30, 2023

## 1 Machine Learning

Textbook is available @ <https://www.github.com/a-mhamdi/mlpy>

---

Textbook is available @ <https://www.github.com/a-mhamdi/mlpy>

### 1.1 Multiple Linear Regression

Multiple linear regression is a type of regression analysis in which there are multiple independent variables that have an effect on the dependent variable. In multiple linear regression, the goal is to find the linear equation that best explains the relationship between the outcome and the features in  $X$ .

The equation takes the form:

$$y = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_{m-1} x_{m-1}$$

where  $y$  is the dependent variable,  $x_1, x_2, \dots, x_{m-1}$  are the independent variables, and  $\theta_0, \theta_1, \theta_2, \dots, \theta_{m-1}$  are the coefficients that represent the influence of each variable on the output  $y$ . The coefficients are estimated using the data, and the resulting equation can be used later to make predictions on new data.

#### 1.1.1 Importing the libraries

```
[1]: import numpy as np
import pandas as pd
from matplotlib import pyplot as plt
```

```
[2]: np.set_printoptions(precision=3)
```

```
[3]: # Show plots in an interactive format, e.g., zooming, saving, etc
%matplotlib inline
```

```
[4]: plt.style.use('ggplot')
```

### 1.1.2 Importing the dataset

```
[5]: df = pd.read_csv('./datasets/50_Startups.csv')
```

```
[6]: df.head()
```

```
[6]:
```

	R&D Spend	Administration	Marketing Spend	State	Profit
0	165349.20	136897.80	471784.10	New York	192261.83
1	162597.70	151377.59	443898.53	California	191792.06
2	153441.51	101145.55	407934.54	Florida	191050.39
3	144372.41	118671.85	383199.62	New York	182901.99
4	142107.34	91391.77	366168.42	Florida	166187.94

```
[7]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50 entries, 0 to 49
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   R&D Spend              50 non-null    float64
1   Administration         50 non-null    float64
2   Marketing Spend        50 non-null    float64
3   State                  50 non-null    object
4   Profit                 50 non-null    float64
dtypes: float64(4), object(1)
memory usage: 2.1+ KB
```

```
[8]: df.describe()
```

```
[8]:
```

	R&D Spend	Administration	Marketing Spend	Profit
count	50.000000	50.000000	50.000000	50.000000
mean	73721.615600	121344.639600	211025.097800	112012.639200
std	45902.256482	28017.802755	122290.310726	40306.180338
min	0.000000	51283.140000	0.000000	14681.400000
25%	39936.370000	103730.875000	129300.132500	90138.902500
50%	73051.080000	122699.795000	212716.240000	107978.190000
75%	101602.800000	144842.180000	299469.085000	139765.977500
max	165349.200000	182645.560000	471784.100000	192261.830000

Extract features  $X$  and target  $y$  from the dataset. **Profit** is the dependant variable.

```
[9]: X = df.iloc[:, :-1]
     y = df.iloc[:, -1]
```

Check the first five observations within  $X$

```
[10]: X.head()
```

```
[10]: R&D Spend Administration Marketing Spend State
0 165349.20 136897.80 471784.10 New York
1 162597.70 151377.59 443898.53 California
2 153441.51 101145.55 407934.54 Florida
3 144372.41 118671.85 383199.62 New York
4 142107.34 91391.77 366168.42 Florida
```

```
[11]: X = X.values
type(X)
```

```
[11]: numpy.ndarray
```

Check the corresponding first five values from **Profit** column.

```
[12]: y.head()
```

```
[12]: 0 192261.83
1 191792.06
2 191050.39
3 182901.99
4 166187.94
Name: Profit, dtype: float64
```

```
[13]: y = y.values
type(y)
```

```
[13]: numpy.ndarray
```

### 1.1.3 Encoding categorical data

```
[14]: from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder
```

```
[15]: ct = ColumnTransformer(transformers=[('encoder', OneHotEncoder(), [3])],
↳remainder='passthrough')
X = np.array(ct.fit_transform(X))
```

```
[16]: print(X[:5])
```

```
[[0.0 0.0 1.0 165349.2 136897.8 471784.1]
 [1.0 0.0 0.0 162597.7 151377.59 443898.53]
 [0.0 1.0 0.0 153441.51 101145.55 407934.54]
 [0.0 0.0 1.0 144372.41 118671.85 383199.62]
 [0.0 1.0 0.0 142107.34 91391.77 366168.42]]
```

#### 1.1.4 Splitting the dataset into training set and test set

```
[17]: from sklearn.model_selection import train_test_split
```

```
[18]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,  
↳ random_state=123)
```

#### 1.1.5 Training the multiple linear regression model on the training set

```
[19]: from sklearn.linear_model import LinearRegression
```

This code will create a linear regression model that fits a line to the training data, in order to make future predictions on the test data.

```
[20]: lr = LinearRegression()  
lr.fit(X_train, y_train)
```

```
[20]: LinearRegression()
```

```
[21]: theta = lr.coef_  
theta
```

```
[21]: array([-1.455e+02, -4.153e+02,  5.607e+02,  7.753e-01, -1.645e-02,  
          3.627e-02])
```

```
[22]: b = lr.intercept_  
b
```

```
[22]: 48661.699896543345
```

Consider the sample `tst` as follows:

```
[23]: tst = np.array([1, 0, 0, 15e+3, 10e+2, 5e+6])
```

Predict the outcome if `tst` is the input.

```
[24]: tst = np.array([1, 0, 0, 15e+3, 10e+2, 5e+6])  
pred = theta @ tst + b  
print('%.3f' % pred)
```

```
241495.528
```

By calling our `lr`, we get the same result:

```
[25]: lr.predict(tst.reshape(1, -1))
```

```
[25]: array([241495.528])
```



If we don't want to do the encoding of state feature by ourselves, we can invoke the previous `ct` object.

```
[26]: tst_new = [[15e+3, 10e+2, 5e+6, 'California']]
      arr = np.array(ct.transform(tst_new))
      arr
```

```
[26]: array([[1.0, 0.0, 0.0, 15000.0, 1000.0, 5000000.0]], dtype=object)
```

```
[27]: lr.predict(arr)
```


```
[27]: array([241495.528])
```

### 1.1.6 Evaluation and Visualization

Make predictions using the  $X$  test set and visualize the results

```
[28]: y_pred = lr.predict(X_test)
```

```
[29]: plt.scatter(y_test, y_pred, c='red') # INSIGHT ON THE CORRELATION BETWEEN
      ↪ Y_TEST AND Y_PRED
      plt.grid()
```



Multiple linear regression can be used to understand the relationship between multiple independent variables and a single dependent variable, and can be used to make predictions about the dependent variable given new data. However, it's important to note that the independent variables must be linearly related to the dependent variable in order for multiple linear regression to behave appropriately. If the relationship is non-linear, we need to use a different type of regression analysis such as polynomial regression.

### 3 | $k$ -NN for Classification

<b>Student's name</b>	.....	.....	.....
	.....	.....	.....
	.....	.....	.....
<b>Score</b> /20	.....	.....	.....

#### Detailed Credits

<b>Anticipation (4 points)</b>	.....	.....	.....
<b>Management (2 points)</b>	.....	.....	.....
<b>Testing (7 points)</b>	.....	.....	.....
<b>Data Logging (3 points)</b>	.....	.....	.....
<b>Interpretation (4 points)</b>	.....	.....	.....

#### Motivations

- ★  $k$ -nearest neighbors ( $k$ -NN) is a simple and effective classification algorithm that is easy to understand and implement. It is based on the idea of using the class labels of the "nearest neighbors" to predict the class label of a new data point.
- ★  $k$ -NN is a "lazy learner" that does not make any assumptions about the underlying data distribution, which makes it a good choice for working with complex or non-linear data. It is also robust to noise and can handle missing data. As a result,  $k$ -NN is often used as a baseline method for comparison with more advanced classification algorithms.



The notebook is available at <https://github.com/a-mhamdi/mlpy/> → Codes → Python →  $k$ -nearest-neighbors.ipynb

# k-nearest-neighbors

November 30, 2023

## 1 Machine Learning

Textbook is available @ <https://www.github.com/a-mhamdi/mlpy>

---

### 1.1 K-Nearest Neighbors (K-NN)

$k$ -nearest neighbors ( $k$ -NN) is a type of instance-based learning, a method of supervised machine learning. It is used for classification and regression tasks.

In  $k$ -NN, the algorithm is given a labeled training dataset and a set of test data. To make a prediction for a test instance, the algorithm looks at the  $k$  nearest neighbors in the training dataset, based on the distance between the test instance and the training instances. The prediction is then made based on the majority class among the  $k$  nearest neighbors. For classification tasks, the prediction is the class with the most neighbors. For regression tasks, the prediction is the mean or median of the values of the  $k$  nearest neighbors.

#### 1.1.1 Importing the libraries

```
[1]: import pandas as pd
```

#### 1.1.2 Importing the dataset

```
[2]: df = pd.read_csv('./datasets/Social_Network_Ads.csv')
df.head()
```

```
[2]:
```

	Age	EstimatedSalary	Purchased
0	19	19000	0
1	35	20000	0
2	26	43000	0
3	27	57000	0
4	19	76000	0

```
[3]: X = df.iloc[:, :-1].values
y = df.iloc[:, -1].values
```

### 1.1.3 Splitting the dataset into the Training set and Test set

```
[4]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,
random_state=123)
```

### 1.1.4 Feature Scaling

$k$ -NN is sensitive to the scale of the features, and it may not perform well if the features have very different scales.

```
[5]: from sklearn.preprocessing import StandardScaler
```

In order to avoid *information leakage*, it is highly important to keep in mind that only the **transform** method has to be applied on the **X\_test**. ( $\mu$ ,  $\sigma$ ) are of **X\_train** set.

```
[6]: sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

### 1.1.5 Training the k-NN model on the training set

```
[7]: from sklearn.neighbors import KNeighborsClassifier
```

```
[8]: clf = KNeighborsClassifier(n_neighbors=5, metric='minkowski', p=2)
```

```
[9]: clf.fit(X_train, y_train)
```

```
[9]: KNeighborsClassifier()
```

### 1.1.6 Predicting a new result

```
[10]: clf.predict(sc.transform([[30,87000]]))
```

```
[10]: array([0])
```

### 1.1.7 Predicting the test set results

```
[11]: y_pred = clf.predict(X_test)
```

### 1.1.8 Displaying the Confusion Matrix

```
[12]: from sklearn.metrics import confusion_matrix
```


```
[13]: cm = confusion_matrix(y_test, y_pred, labels=clf.classes_)
```

```
[14]: cm
```

```
[14]: array([[54,  9],  
        [ 4, 33]])
```

```
[15]: from sklearn.metrics import ConfusionMatrixDisplay
```

```
[16]: ConfusionMatrixDisplay(cm, display_labels=clf.classes_).plot();
```



```
[17]: from sklearn.metrics import accuracy_score
```

```
[18]: print(f'Accuracy = {accuracy_score(y_test, y_pred):.2f}')
```

Accuracy = 0.87

<https://pandas.pydata.org/pandas-docs/stable/reference/api/pandas.crosstab.html>

```
[19]: pd.crosstab(y_test, y_pred, rownames=['Expected'], colnames=['Predicted'],  
               margins=True)
```

```
[19]: Predicted    0    1  All  
      Expected
```

0	54	9	63
1	4	33	37
All	58	42	100

```
[20]: from sklearn.metrics import classification_report
```

```
[21]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.93	0.86	0.89	63
1	0.79	0.89	0.84	37
accuracy			0.87	100
macro avg	0.86	0.87	0.86	100
weighted avg	0.88	0.87	0.87	100

$k$ -NN is a simple and effective method for classification and regression tasks, and it is easy to understand and implement. However, it can be computationally expensive to find the  $k$  nearest neighbors for each test instance, especially for large datasets.

## 4 | K-Means for Clustering

<b>Student's name</b>	.....	.....	.....
	.....	.....	.....
	.....	.....	.....
<b>Score</b> /20	.....	.....	.....

### Detailed Credits

<b>Anticipation (4 points)</b>	.....	.....	.....
<b>Management (2 points)</b>	.....	.....	.....
<b>Testing (7 points)</b>	.....	.....	.....
<b>Data Logging (3 points)</b>	.....	.....	.....
<b>Interpretation (4 points)</b>	.....	.....	.....

#### Motivations

- ★ K-means clustering is a widely used method for partitioning a dataset into a set of clusters, where each cluster consists of data points that are similar to each other. This can be useful for a variety of applications, including data compression, anomaly detection, and customer segmentation.
- ★ It can also help to identify outliers and anomalies in the data, which can be useful for identifying errors or identifying new opportunities for analysis.



The notebook is available at <https://github.com/a-mhamdi/mlpy/> → Codes → Python → k-means-clustering.ipynb



# k-means-clustering

November 30, 2023

## 1 Machine Learning

Textbook is available @ <https://www.github.com/a-mhamdi/mlpy>

---

In unsupervised learning, the algorithm is given a dataset and is asked to learn the underlying structure of the data. The goal is to find patterns or relationships in the data that can be used to group the data points into clusters or to reduce the dimensionality of the data.

Some examples of unsupervised learning algorithms include: 1.  $K$ -means clustering; 1. Principal Component Analysis (PCA); and 1. Autoencoders.

These algorithms can be used for tasks such as image compression, anomaly detection, and customer segmentation.

### 1.1 K-Means Clustering

$K$ -means clustering is a method of unsupervised machine learning used to partition a dataset into  $k$  clusters, where  $k$  is a user-specified number. The goal of  $K$ -means clustering is to minimize the sum of squared distances between the points in each cluster and its centroid.

#### 1.1.1 Importing the libraries

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
```

```
[2]: # Show plots in an interactive format, e.g., zooming, saving, etc
%matplotlib inline
```

```
[3]: plt.style.use('ggplot')
```

#### 1.1.2 Importing the dataset

```
[4]: df = pd.read_csv('./datasets/Mall_Customers.csv')
```

```
[5]: df.head()
```

```
[5]:
```

	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	Male	19	15	39
1	2	Male	21	15	81
2	3	Female	20	16	6
3	4	Female	23	16	77
4	5	Female	31	17	40

```
[6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   CustomerID            200 non-null   int64
1   Gender                200 non-null   object
2   Age                   200 non-null   int64
3   Annual Income (k$)    200 non-null   int64
4   Spending Score (1-100) 200 non-null   int64
dtypes: int64(4), object(1)
memory usage: 7.9+ KB
```

```
[7]: df.describe()
```

```
[7]:
```

	CustomerID	Age	Annual Income (k\$)	Spending Score (1-100)
count	200.000000	200.000000	200.000000	200.000000
mean	100.500000	38.850000	60.560000	50.200000
std	57.879185	13.969007	26.264721	25.823522
min	1.000000	18.000000	15.000000	1.000000
25%	50.750000	28.750000	41.500000	34.750000
50%	100.500000	36.000000	61.500000	50.000000
75%	150.250000	49.000000	78.000000	73.000000
max	200.000000	70.000000	137.000000	99.000000

```
[8]: df.rename(columns={'Annual Income (k$)': 'Annual Income', 'Spending Score_
↳ (1-100)': 'Spending Score'}, inplace=True)
```

```
[9]: X = df.drop(columns=['CustomerID', 'Age', 'Gender']).values
X[:10, :]
```

```
[9]: array([[15, 39],
           [15, 81],
           [16, 6],
           [16, 77],
           [17, 40],
           [17, 76],
           [18, 6],
```

```
[18, 94],  
[19, 3],  
[19, 72]])
```

Import KMeans class

```
[10]: from sklearn.cluster import KMeans
```

**OPTIONAL: IF NOT FAMILIAR WITH KMEANS, FEEL FREE TO SKIP THE FOLLOWING CELL**

---

Using the elbow method to find the optimal number of clusters

```
[11]: """  
wcss = []  
  
for i in range(1, 11):  
    kmeans = KMeans(n_clusters=i,  
                    init='k-means++', # Init method  
                    random_state=123) # Random seed for reproducibility  
    kmeans.fit(X)  
    wcss.append(kmeans.inertia_)  
  
plt.plot(range(1, 11), wcss)  
plt.title('The Elbow Method')  
plt.xlabel('Number of clusters')  
plt.ylabel('WCSS')  
plt.grid()  
"""
```

```
[11]: "\nwcss = []\n\nfor i in range(1, 11):\n    kmeans = KMeans(n_clusters=i,\n                    init='k-means++', # Init method\n                    random_state=123) # Random\nseed for reproducibility\n    kmeans.fit(X)\n    wcss.append(kmeans.inertia_)\n\nplt.plot(range(1, 11), wcss)\nplt.title('The\nElbow Method')\nplt.xlabel('Number of\nclusters')\nplt.ylabel('WCSS')\nplt.grid()\n"
```

---

### 1.1.3 Training the K-Means model on the dataset

This code will create a *K*-means model with 5 clusters and fit it to the data. It will then make predictions about which cluster each data point belongs to

```
[12]: kmeans = KMeans(n_clusters=5, init='k-means++', random_state=123)  
y_pred = kmeans.fit_predict(X)
```

```
/home/mhamdi/MEGA/git-repos/AI-ML-  
DL/mlpy/Codes/Python/pvenv/lib/python3.10/site-  
packages/sklearn/cluster/_kmeans.py:870: FutureWarning: The default value of  
`n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init`  
explicitly to suppress the warning  
warnings.warn(  


```

```
[13]: centers = kmeans.cluster_centers_  
centers
```

```
[13]: array([[55.2962963 , 49.51851852],  
            [86.53846154, 82.12820513],  
            [88.2       , 17.11428571],  
            [26.30434783, 20.91304348],  
            [25.72727273, 79.36363636]])
```

#### 1.1.4 Visualizing the clusters

```
[14]: fig, ax = plt.subplots()  
scatter = ax.scatter(X[:, 0], X[:, 1], c=y_pred, s=100)  
legend = ax.legend(*scatter.legend_elements(), title='Clusters')  
ax.add_artist(legend)  
ax.scatter(centers[:, 0], centers[:, 1], c='red', s=200)  
ax.set_title('Clusters of customers')  
ax.set_xlabel('Annual Income')  
ax.set_ylabel('Spending Score')  
ax.grid()
```



Unsupervised learning can be useful when there is no labeled training data available, or when the goal is to discover patterns or relationships in the data rather than to make predictions. However, it can be more difficult to evaluate the performance of unsupervised learning algorithms, as there is no ground truth to compare the predictions to.

*K*-means clustering is a fast and efficient method for clustering large datasets, and is often used as a baseline method for comparison with other clustering algorithms. However, it can be sensitive to the initial selection of centroids, and may not always find the optimal clusters if the data is not well-separated or has a non-convex shape. It is also limited to spherical clusters and may not work well for clusters with more complex shapes.

## 5 | Binary Classifier using ANN

<b>Student's name</b>	.....	.....	.....
	.....	.....	.....
	.....	.....	.....
<b>Score</b> <b>/20</b>	.....	.....	.....

### Detailed Credits

<b>Anticipation (4 points)</b>	.....	.....	.....
<b>Management (2 points)</b>	.....	.....	.....
<b>Testing (7 points)</b>	.....	.....	.....
<b>Data Logging (3 points)</b>	.....	.....	.....
<b>Interpretation (4 points)</b>	.....	.....	.....

#### Motivations

- ★ Artificial neural networks (ANNs) are a powerful tool for binary classification tasks, which involve predicting a binary outcome (e.g., “yes” or “no”) based on input data. ANNs are able to learn complex relationships between the input data and the output labels, which makes them well-suited for tasks with a large number of features or a complex underlying structure.
- ★ ANNs are highly flexible and can be trained on a wide range of data types, including continuous and categorical variables. They can also handle missing values and handle large amounts of data efficiently. This makes them a good choice for tasks where the data is noisy or high-dimensional.



The notebook is available at <https://github.com/a-mhamdi/mlpy/> → Codes → Python → artificial-neural-network.ipynb

# artificial-neural-network

November 30, 2023

## 1 Machine Learning

Textbook is available @ <https://www.github.com/a-mhamdi/mlpy>

---

Artificial neural networks (ANN) are commonly used for classification tasks because they are able to learn complex relationships between the input features and the target class. They are particularly useful when the relationship is non-linear, as they are able to learn and model the inputs-outputs mapping using multiple hidden layers of interconnected neurons.

ANN are also able to handle large amounts of data and can learn from it without being explicitly programmed with a set of rules or a decision tree. This allows them to be very flexible and adaptable, and makes them well-suited for tasks that are difficult to define using traditional programming techniques.

### 1.1 Binary Classification using ANN

There are several advantages to using neural networks for classification tasks:

1. They are able to learn complex relationships between the input features and the target class;
2. They are able to handle large amounts of data;
3. They can learn from unstructured data;
4. They are flexible and adaptable;
5. They can be trained to perform well on a wide range of classification tasks.

#### 1.1.1 Importing the libraries

```
[1]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
```

```
[2]: np.set_printoptions(precision=2)
```

#### 1.1.2 Importing the dataset

```
[3]: df = pd.read_csv("./datasets/Churn_Modelling.csv")
```

```
[4]: df = df.dropna(how="any", axis=0)
```

```
[5]: df.head()
```

```
[5]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	\
0	1	15634602	Hargrave	619	France	Female	42	
1	2	15647311	Hill	608	Spain	Female	41	
2	3	15619304	Onio	502	France	Female	42	
3	4	15701354	Boni	699	France	Female	39	
4	5	15737888	Mitchell	850	Spain	Female	43	

	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
0	2	0.00	1	1	1	
1	1	83807.86	1	0	1	
2	8	159660.80	3	1	0	
3	1	0.00	2	0	0	
4	2	125510.82	1	1	1	

	EstimatedSalary	Exited
0	101348.88	1
1	112542.58	0
2	113931.57	1
3	93826.63	0
4	79084.10	0

```
[6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   RowNumber              10000 non-null  int64
1   CustomerId             10000 non-null  int64
2   Surname                 10000 non-null  object
3   CreditScore             10000 non-null  int64
4   Geography               10000 non-null  object
5   Gender                  10000 non-null  object
6   Age                     10000 non-null  int64
7   Tenure                  10000 non-null  int64
8   Balance                 10000 non-null  float64
9   NumOfProducts           10000 non-null  int64
10  HasCrCard               10000 non-null  int64
11  IsActiveMember          10000 non-null  int64
12  EstimatedSalary          10000 non-null  float64
13  Exited                   10000 non-null  int64
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB
```



```
[7]: df.describe()
```

```
[7]:
```

	RowNumber	CustomerId	CreditScore	Age	Tenure \
count	10000.00000	1.000000e+04	10000.000000	10000.000000	10000.000000
mean	5000.50000	1.569094e+07	650.528800	38.921800	5.012800
std	2886.89568	7.193619e+04	96.653299	10.487806	2.892174
min	1.00000	1.556570e+07	350.000000	18.000000	0.000000
25%	2500.75000	1.562853e+07	584.000000	32.000000	3.000000
50%	5000.50000	1.569074e+07	652.000000	37.000000	5.000000
75%	7500.25000	1.575323e+07	718.000000	44.000000	7.000000
max	10000.00000	1.581569e+07	850.000000	92.000000	10.000000

	Balance	NumOfProducts	HasCrCard	IsActiveMember \
count	10000.000000	10000.000000	10000.000000	10000.000000
mean	76485.889288	1.530200	0.70550	0.515100
std	62397.405202	0.581654	0.45584	0.499797
min	0.000000	1.000000	0.00000	0.000000
25%	0.000000	1.000000	0.00000	0.000000
50%	97198.540000	1.000000	1.00000	1.000000
75%	127644.240000	2.000000	1.00000	1.000000
max	250898.090000	4.000000	1.00000	1.000000

	EstimatedSalary	Exited
count	10000.000000	10000.000000
mean	100090.239881	0.203700
std	57510.492818	0.402769
min	11.580000	0.000000
25%	51002.110000	0.000000
50%	100193.915000	0.000000
75%	149388.247500	0.000000
max	199992.480000	1.000000

```
[8]: X = df.iloc[:, 3:-1].values
     y = df.iloc[:, -1].values
```

### 1.1.3 Data preprocessing

```
[9]: from sklearn.preprocessing import LabelEncoder, OneHotEncoder
```

```
[10]: le = LabelEncoder()
     ohe = OneHotEncoder()
```

```
[11]: X[:, 2] = le.fit_transform(X[:, 2])
```

```
[12]: from sklearn.compose import ColumnTransformer
```

```
[13]: ct = ColumnTransformer([("ohe", ohe, [1])], remainder='passthrough')
X = np.array(ct.fit_transform(X))
```

```
[14]: X[:5, :]
```

```
[14]: array([[1.0, 0.0, 0.0, 619, 0, 42, 2, 0.0, 1, 1, 1, 101348.88],
          [0.0, 0.0, 1.0, 608, 0, 41, 1, 83807.86, 1, 0, 1, 112542.58],
          [1.0, 0.0, 0.0, 502, 0, 42, 8, 159660.8, 3, 1, 0, 113931.57],
          [1.0, 0.0, 0.0, 699, 0, 39, 1, 0.0, 2, 0, 0, 93826.63],
          [0.0, 0.0, 1.0, 850, 0, 43, 2, 125510.82, 1, 1, 1, 79084.1]],
        dtype=object)
```

```
[15]: X = np.asarray(X, dtype=np.float64)
```

```
[16]: X[:5, :]
```

```
[16]: array([[1.00e+00, 0.00e+00, 0.00e+00, 6.19e+02, 0.00e+00, 4.20e+01,
          2.00e+00, 0.00e+00, 1.00e+00, 1.00e+00, 1.00e+00, 1.01e+05],
          [0.00e+00, 0.00e+00, 1.00e+00, 6.08e+02, 0.00e+00, 4.10e+01,
          1.00e+00, 8.38e+04, 1.00e+00, 0.00e+00, 1.00e+00, 1.13e+05],
          [1.00e+00, 0.00e+00, 0.00e+00, 5.02e+02, 0.00e+00, 4.20e+01,
          8.00e+00, 1.60e+05, 3.00e+00, 1.00e+00, 0.00e+00, 1.14e+05],
          [1.00e+00, 0.00e+00, 0.00e+00, 6.99e+02, 0.00e+00, 3.90e+01,
          1.00e+00, 0.00e+00, 2.00e+00, 0.00e+00, 0.00e+00, 9.38e+04],
          [0.00e+00, 0.00e+00, 1.00e+00, 8.50e+02, 0.00e+00, 4.30e+01,
          2.00e+00, 1.26e+05, 1.00e+00, 1.00e+00, 1.00e+00, 7.91e+04]])
```

```
[17]: from sklearn.model_selection import train_test_split
```

```
[18]: X_train, X_test, y_train, y_test = train_test_split(X, y, train_size=.8,
        random_state=123)
```

```
[19]: from sklearn.preprocessing import MinMaxScaler
```

```
[20]: sc = MinMaxScaler()
```

```
[21]: X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

```
[22]: print(X_train[:5, :])
```

```
[[1.  0.  0.  0.75 1.  0.16 1.  0.  0.33 1.  1.  0.27]
 [1.  0.  0.  0.51 0.  0.28 1.  0.  0.67 1.  0.  0.66]
 [0.  0.  1.  0.87 1.  0.23 0.3  0.  0.33 0.  0.  0.41]
 [1.  0.  0.  0.69 1.  0.3  0.9  0.  0.33 1.  0.  0.2 ]
 [1.  0.  0.  0.71 1.  0.2  0.3  0.58 0.33 1.  0.  0.57]]
```

#### 1.1.4 Build the classifier clf

```
[23]: from keras.models import Sequential
      from keras.layers import Dense
```

```
2023-11-30 01:23:04.966677: I tensorflow/core/platform/cpu_feature_guard.cc:193]
This TensorFlow binary is optimized with oneAPI Deep Neural Network Library
(oneDNN) to use the following CPU instructions in performance-critical
operations:  AVX2 FMA
To enable them in other operations, rebuild TensorFlow with the appropriate
compiler flags.
2023-11-30 01:23:05.536785: W
tensorflow/compiler/xla/stream_executor/platform/default/dso_loader.cc:64] Could
not load dynamic library 'libcudart.so.11.0'; dlerror: libcudart.so.11.0: cannot
open shared object file: No such file or directory
2023-11-30 01:23:05.536804: I
tensorflow/compiler/xla/stream_executor/cuda/cudart_stub.cc:29] Ignore above
cudart dlerror if you do not have a GPU set up on your machine.
2023-11-30 01:23:06.697552: W
tensorflow/compiler/xla/stream_executor/platform/default/dso_loader.cc:64] Could
not load dynamic library 'libnvinfer.so.7'; dlerror: libnvinfer.so.7: cannot
open shared object file: No such file or directory
2023-11-30 01:23:06.697725: W
tensorflow/compiler/xla/stream_executor/platform/default/dso_loader.cc:64] Could
not load dynamic library 'libnvinfer_plugin.so.7'; dlerror:
libnvinfer_plugin.so.7: cannot open shared object file: No such file or
directory
2023-11-30 01:23:06.697739: W
tensorflow/compiler/tf2tensorrt/utils/py_utils.cc:38] TF-TRT Warning: Cannot
dlopen some TensorRT libraries. If you would like to use Nvidia GPU with
TensorRT, please make sure the missing libraries mentioned above are installed
properly.
```

```
[24]: clf = Sequential()
      ndim = X_train.shape[1]
      clf.add(Dense(units=8, activation="relu", input_dim=ndim))
      clf.add(Dense(units=4, activation="relu"))
      clf.add(Dense(units=4, activation="relu"))
      clf.add(Dense(units=1, activation="sigmoid"))
```

```
2023-11-30 01:23:07.562116: W
tensorflow/compiler/xla/stream_executor/platform/default/dso_loader.cc:64] Could
not load dynamic library 'libcuda.so.1'; dlerror: libcuda.so.1: cannot open
shared object file: No such file or directory
2023-11-30 01:23:07.562177: W
tensorflow/compiler/xla/stream_executor/cuda/cuda_driver.cc:265] failed call to
cuInit: UNKNOWN ERROR (303)
2023-11-30 01:23:07.562196: I
```

```
tensorflow/compiler/xla/stream_executor/cuda/cuda_diagnostics.cc:156] kernel
driver does not appear to be running on this host (e590):
/proc/driver/nvidia/version does not exist
2023-11-30 01:23:07.562409: I tensorflow/core/platform/cpu_feature_guard.cc:193]
This TensorFlow binary is optimized with oneAPI Deep Neural Network Library
(oneDNN) to use the following CPU instructions in performance-critical
operations: AVX2 FMA
To enable them in other operations, rebuild TensorFlow with the appropriate
compiler flags.
```

### 1.1.5 Insights about clf

```
[25]: clf.summary()
```

```
Model: "sequential"
```

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 8)	104
dense_1 (Dense)	(None, 4)	36
dense_2 (Dense)	(None, 4)	20
dense_3 (Dense)	(None, 1)	5

```

Total params: 165
Trainable params: 165
Non-trainable params: 0

```

### 1.1.6 Compile clf

```
[26]: clf.compile(optimizer="adam",
                  loss="binary_crossentropy",
                  metrics=['Accuracy', 'Precision', 'Recall'])
```

### 1.1.7 Train and evaluate clf

```
[27]: clf.fit(X_train, y_train, batch_size=16, epochs=32);
```

```
Epoch 1/32
500/500 [=====] - 1s 1ms/step - loss: 0.5106 -
Accuracy: 0.7904 - precision: 0.3533 - recall: 0.0400
Epoch 2/32
500/500 [=====] - 1s 1ms/step - loss: 0.4553 -
Accuracy: 0.7971 - precision: 0.0000e+00 - recall: 0.0000e+00
```

Epoch 3/32  
500/500 [=====] - 1s 1ms/step - loss: 0.4392 -  
Accuracy: 0.7971 - precision: 0.0000e+00 - recall: 0.0000e+00

Epoch 4/32  
500/500 [=====] - 1s 1ms/step - loss: 0.4218 -  
Accuracy: 0.7983 - precision: 0.8000 - recall: 0.0074

Epoch 5/32  
500/500 [=====] - 1s 1ms/step - loss: 0.4056 -  
Accuracy: 0.8145 - precision: 0.6980 - recall: 0.1510

Epoch 6/32  
500/500 [=====] - 1s 1ms/step - loss: 0.3941 -  
Accuracy: 0.8254 - precision: 0.6733 - recall: 0.2705

Epoch 7/32  
500/500 [=====] - 1s 1ms/step - loss: 0.3883 -  
Accuracy: 0.8319 - precision: 0.6853 - recall: 0.3167

Epoch 8/32  
500/500 [=====] - 1s 1ms/step - loss: 0.3840 -  
Accuracy: 0.8364 - precision: 0.6817 - recall: 0.3629

Epoch 9/32  
500/500 [=====] - 1s 1ms/step - loss: 0.3799 -  
Accuracy: 0.8364 - precision: 0.6718 - recall: 0.3783

Epoch 10/32  
500/500 [=====] - 1s 1ms/step - loss: 0.3779 -  
Accuracy: 0.8400 - precision: 0.6850 - recall: 0.3913

Epoch 11/32  
500/500 [=====] - 1s 1ms/step - loss: 0.3753 -  
Accuracy: 0.8440 - precision: 0.7045 - recall: 0.3980

Epoch 12/32  
500/500 [=====] - 1s 1ms/step - loss: 0.3728 -  
Accuracy: 0.8465 - precision: 0.7112 - recall: 0.4097

Epoch 13/32  
500/500 [=====] - 1s 1ms/step - loss: 0.3711 -  
Accuracy: 0.8470 - precision: 0.7157 - recall: 0.4079

Epoch 14/32  
500/500 [=====] - 1s 1ms/step - loss: 0.3700 -  
Accuracy: 0.8489 - precision: 0.7216 - recall: 0.4153

Epoch 15/32  
500/500 [=====] - 1s 1ms/step - loss: 0.3688 -  
Accuracy: 0.8497 - precision: 0.7223 - recall: 0.4214

Epoch 16/32  
500/500 [=====] - 1s 1ms/step - loss: 0.3673 -  
Accuracy: 0.8499 - precision: 0.7279 - recall: 0.4153

Epoch 17/32  
500/500 [=====] - 1s 1ms/step - loss: 0.3673 -  
Accuracy: 0.8503 - precision: 0.7216 - recall: 0.4264

Epoch 18/32  
500/500 [=====] - 1s 1ms/step - loss: 0.3653 -  
Accuracy: 0.8528 - precision: 0.7350 - recall: 0.4288

```

Epoch 19/32
500/500 [=====] - 1s 1ms/step - loss: 0.3650 -
Accuracy: 0.8524 - precision: 0.7361 - recall: 0.4245
Epoch 20/32
500/500 [=====] - 1s 1ms/step - loss: 0.3646 -
Accuracy: 0.8529 - precision: 0.7333 - recall: 0.4319
Epoch 21/32
500/500 [=====] - 1s 1ms/step - loss: 0.3629 -
Accuracy: 0.8539 - precision: 0.7441 - recall: 0.4264
Epoch 22/32
500/500 [=====] - 1s 1ms/step - loss: 0.3614 -
Accuracy: 0.8531 - precision: 0.7419 - recall: 0.4233
Epoch 23/32
500/500 [=====] - 1s 1ms/step - loss: 0.3614 -
Accuracy: 0.8551 - precision: 0.7432 - recall: 0.4368
Epoch 24/32
500/500 [=====] - 1s 1ms/step - loss: 0.3603 -
Accuracy: 0.8547 - precision: 0.7444 - recall: 0.4325
Epoch 25/32
500/500 [=====] - 1s 1ms/step - loss: 0.3589 -
Accuracy: 0.8533 - precision: 0.7341 - recall: 0.4338
Epoch 26/32
500/500 [=====] - 1s 1ms/step - loss: 0.3578 -
Accuracy: 0.8545 - precision: 0.7393 - recall: 0.4368
Epoch 27/32
500/500 [=====] - 1s 1ms/step - loss: 0.3568 -
Accuracy: 0.8549 - precision: 0.7416 - recall: 0.4368
Epoch 28/32
500/500 [=====] - 1s 1ms/step - loss: 0.3575 -
Accuracy: 0.8561 - precision: 0.7453 - recall: 0.4418
Epoch 29/32
500/500 [=====] - 1s 1ms/step - loss: 0.3562 -
Accuracy: 0.8558 - precision: 0.7503 - recall: 0.4331
Epoch 30/32
500/500 [=====] - 1s 1ms/step - loss: 0.3549 -
Accuracy: 0.8546 - precision: 0.7478 - recall: 0.4276
Epoch 31/32
500/500 [=====] - 1s 1ms/step - loss: 0.3548 -
Accuracy: 0.8550 - precision: 0.7360 - recall: 0.4449
Epoch 32/32
500/500 [=====] - 1s 1ms/step - loss: 0.3545 -
Accuracy: 0.8558 - precision: 0.7492 - recall: 0.4344

```

```
[28]: from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
```


```
[29]: y_pred = clf.predict(X_test)
      y_pred = ( y_pred > .5 ).astype(int)
```

63/63 [=====] - 0s 773us/step

Print the confusion matrix

```
[30]: cm = confusion_matrix(y_test, y_pred)
```

```
[31]: ConfusionMatrixDisplay(cm).plot();
```



```
[32]: print(y_test.shape)
```

(2000,)

```
[33]: print(y_pred.shape)
y_pred = y_pred.reshape(len(y_pred),)
print(y_pred.shape)
```

(2000, 1)

(2000,)

```
[34]: pd.crosstab(y_test, y_pred, rownames=["Expected"], colnames=["Predicted"],
    margins=True)
```

```
[34]: Predicted      0      1    All
      Expected
      0          1529    57  1586
      1           217   197   414
      All          1746   254  2000
```

```
[35]: y_test = y_test.reshape(len(y_test), 1)
      y_pred = y_pred.reshape(len(y_pred), 1)
      print(np.concatenate((y_test[:10], y_pred[:10]), axis=1))
```

```
[[0 0]
 [0 0]
 [0 0]
 [0 0]
 [0 0]
 [0 0]
 [1 0]
 [1 1]
 [0 0]
 [0 0]]
```

```
[36]: from sklearn.metrics import classification_report
```

```
[37]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.88	0.96	0.92	1586
1	0.78	0.48	0.59	414
accuracy			0.86	2000
macro avg	0.83	0.72	0.75	2000
weighted avg	0.85	0.86	0.85	2000

It is important to note that neural networks can be more computationally intensive to train and may require more data and more time to achieve good performance, compared to some other classification algorithms. Additionally, they can be more difficult to interpret and understand, as they learn patterns in the data through the weights and biases of the network rather than through explicit rules.



The overall scope of this manual is to introduce **Machine Learning**, through some numeric simulations, to the students at the department of **Electrical Engineering**.

The topics discussed in this manuscript are as follow:

- ① Getting started with *Python*
- ② Linear Regression
- ③ Classification
- ④ Clustering
- ⑤ ANN

*Python; Jupyter; NumPy; Matplotlib; scikit-learn; machine learning; linear regression; classification; clustering; deep learning.*